Machine Learning


NLP

What is NLP?

- Natural Language Processing (NLP) refers to AI method of communicating with an intelligent systems using a natural language such as English
- It is the sub-field of AI that is focused on enabling computers to understand and process human language
- The ultimate objective of NLP is to read, decipher, understand, and make sense of the human languages in a manner that is valuable
- Most NLP techniques rely on machine learning to derive meaning from human languages

Uses cased of NLP

- NLP enables the recognition and **prediction of diseases** based on electronic health records and patient's own speech
- Organizations can determine what customers are saying about a service or product by identifying and extracting information in sources like social media (**sentiment analysis**)
- Companies like Yahoo and Google filter and classify your emails with NLP by analyzing text in emails that flow through their servers and **stopping spam** before they even enter your inbox
- To help **identifying fake news**, a system can be developed to determine if a source is accurate or politically biased, detecting if a news source can be trusted or not
- Amazon's Alexa and Apple's Siri are examples of intelligent **voice driven interfaces** that use NLP to respond to vocal prompts
- NLP is also being used in both the search and selection phases of **talent recruitment**
- NLP is particularly booming in the **healthcare industry**

Terminology

- **Phonology**
  - It is study of organizing sound systematically.
- **Morphology**
  - It is a study of construction of words from primitive meaningful units.
- **Morpheme**
  - It is primitive unit of meaning in a language.
- **Syntax**
  - It refers to arranging words to make a sentence. It also involves determining the structural role of words in the sentence and in phrases.
- **Semantics**
  - It is concerned with the meaning of words and how to combine words into meaningful phrases and sentences.

Terminology

- **Pragmatics**
  - It deals with using and understanding sentences in different situations and how the interpretation of the sentence is affected.
- **Discourse**
  - It deals with how the immediately preceding sentence can affect the interpretation of the next sentence.
- **World Knowledge**
  - It includes the general knowledge about the world.

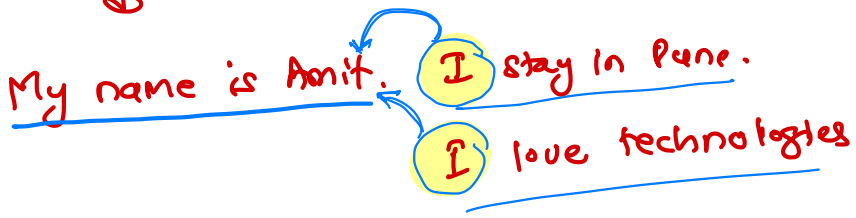

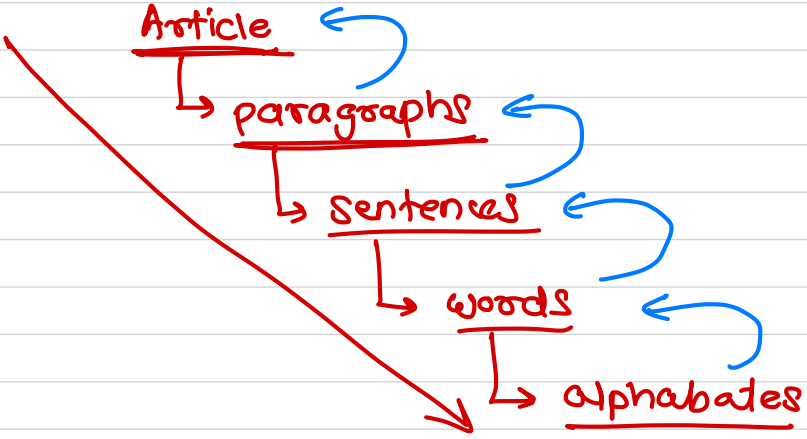

India : place name
Steve : person name

① I like Apple in my breakfast. : fruit

② Apple has again stood on first place ... : company

Steps in NLP

I go to office

Syntax

① Lexical Analysis	It involves identifying and analyzing the structure of words. Lexicon of a language means the collection of words and phrases in a language. Lexical analysis is dividing the whole chunk of txt into paragraphs, sentences, and words → collection of words
② Syntactic Analysis	It involves analysis of words in the sentence for grammar and arranging words in a manner that shows the relationship among the words. The sentence such as "The school goes to boy" is rejected by English syntactic analyzer → grammatically correct sentence
③ Semantic Analysis	It draws the exact meaning or the dictionary meaning from the text. The text is checked for meaningfulness. It is done by mapping syntactic structures and objects in the task domain. ↳ words & their meanings (verb / noun / pronoun)
④ **Disclosure Integration**	The meaning of any sentence depends upon the meaning of the sentence just before it. In addition, it also brings about the meaning of immediately succeeding sentence.
⑤ **Pragmatic Analysis**	During this, what was said is re-interpreted on what it actually meant. It involves deriving those aspects of language which require real world knowledge → apple: fruit / company

① The restaurant is awesome. : good

② The restaurant is not so bad. : good

③ The restaurant is bad. : bad

Creating NLP pipeline

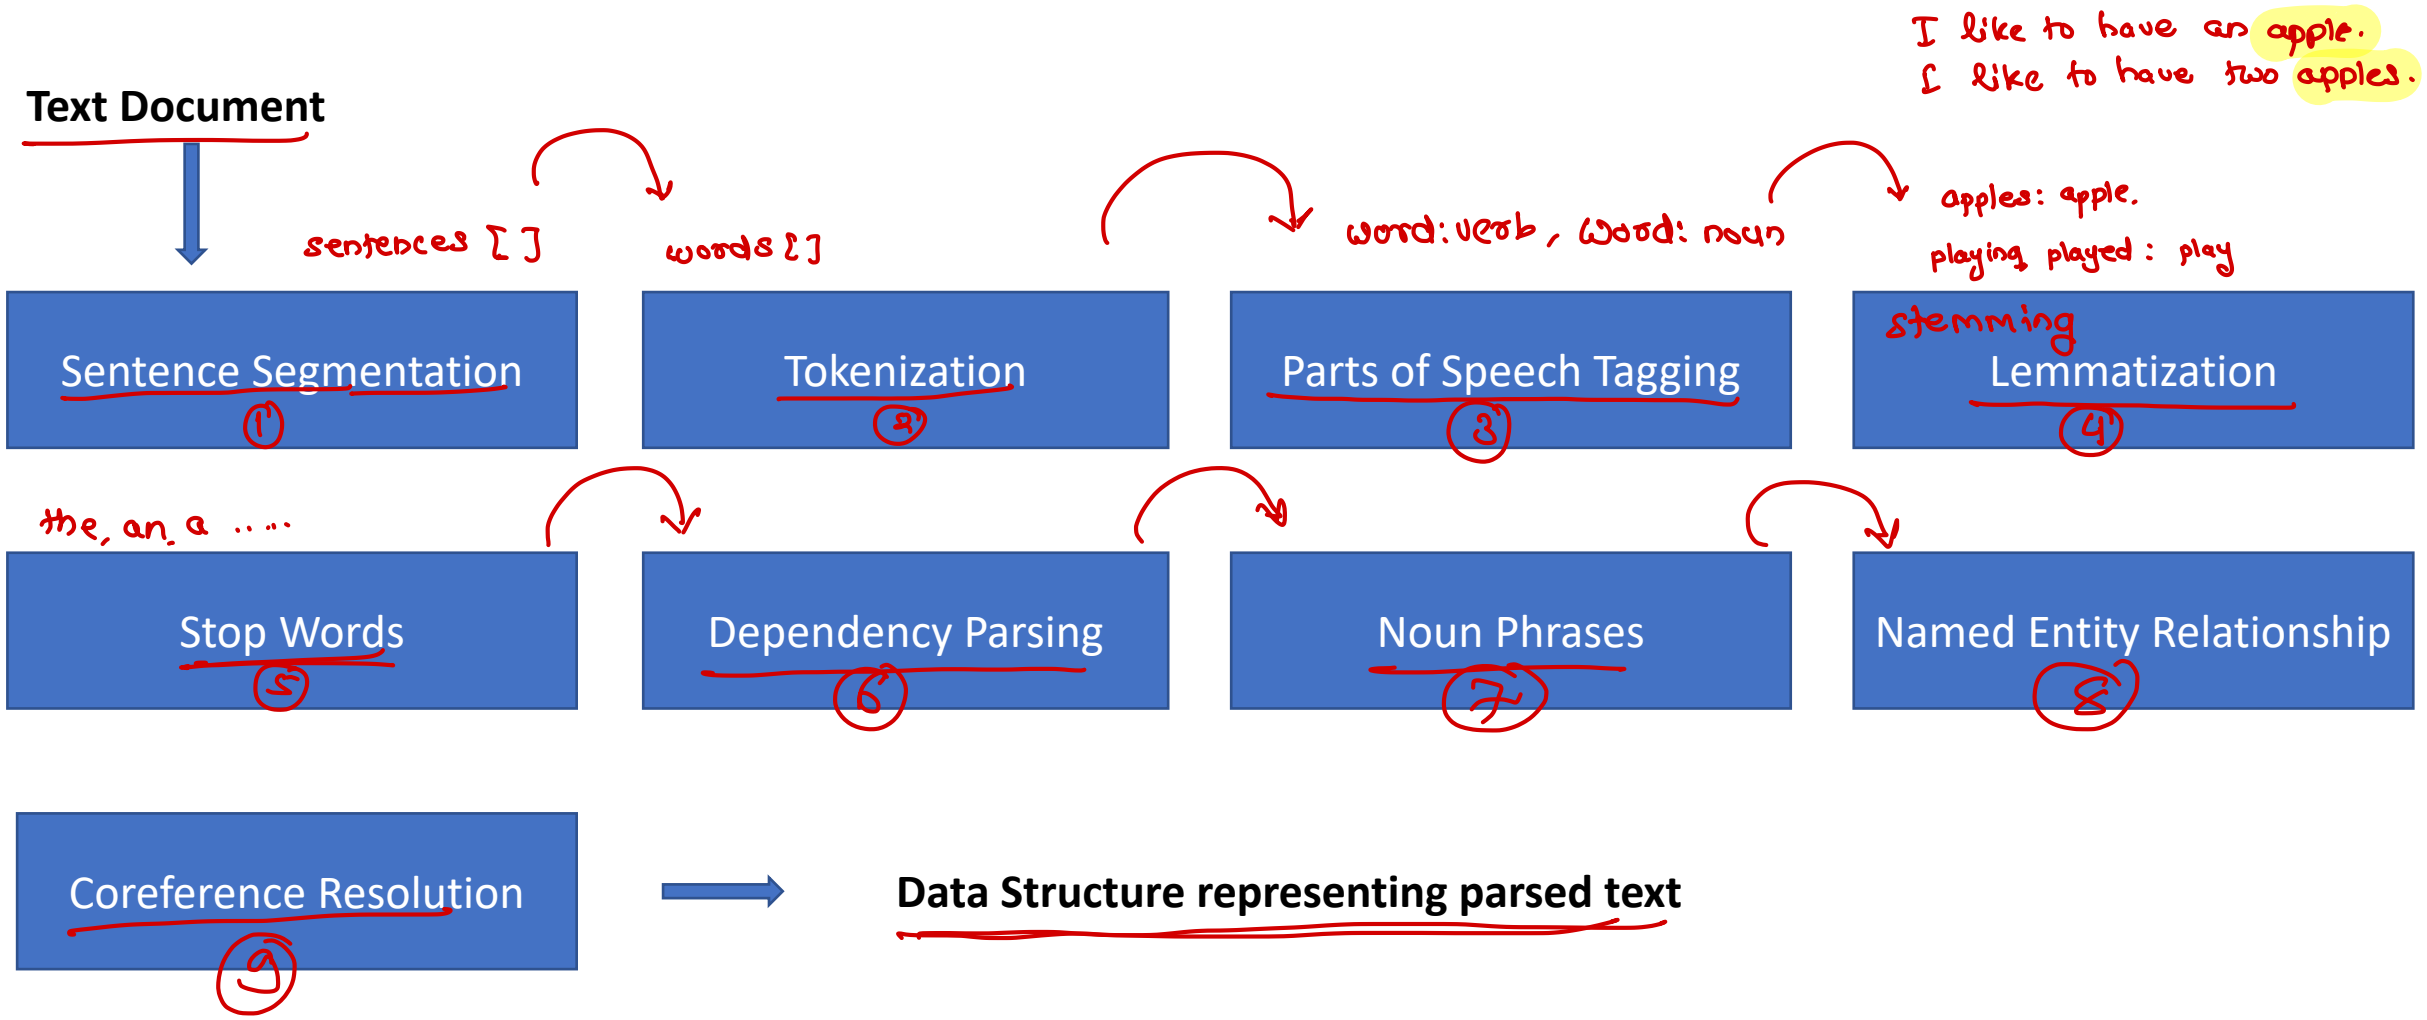

Sentence Segmentation

- The first step in the pipeline is to break the text apart into separate sentences
- We can assume that each sentence in English is a separate thought or idea
- It will be a lot easier to write a program to understand a single sentence than to understand a whole paragraph
- Coding a Sentence Segmentation model can be as simple as splitting apart sentences whenever you see a punctuation mark

Word Tokenization

- Now that we've split our document into sentences, we can process them one at a time
- Tokenization is the process of splitting the sentence into separate words
- Tokenization is easy to do in English. We'll just split apart words whenever there's a space between them. And we'll also treat punctuation marks as separate tokens since punctuation also has meaning.

Predicting Parts of Speech for Each Token

- Next, we'll look at each token and try to guess its part of speech — whether it is a noun, a verb, an adjective and so on
- Knowing the role of each word in the sentence will help us start to figure out what the sentence is talking about
- We can do this by feeding each word (and some extra words around it for context) into a pre-trained part-of-speech classification model
- The part-of-speech model was originally trained by feeding it millions of English sentences with each word's part of speech already tagged and having it learn to replicate that behavior.
- Keep in mind that the model is completely based on statistics — it doesn't actually understand what the words mean in the same way that humans do
- It just knows how to guess a part of speech based on similar sentences and words it has seen before

Text Lemmatization

- In English words appear in different forms. E.g.
  - I ate an apple
  - I ate two apples

} apple

- Both sentences talk about the noun **apple,** but they are using different inflections
- When working with text in a computer, it is helpful to know the base form of each word so that you know that both sentences are talking about the same concept
- Otherwise the strings like apple and apples will look lie two totally different words to computer
- In NLP, we call finding this process *lemmatization* — figuring out the most basic form or *lemma* of each word in the sentence.
- Lemmatization is typically done by having a look-up table of the lemma forms of words based on their part of speech and possibly having some custom rules to handle words that you've never seen before

playing → play

played → play

Identifying Stop Words

- Next, we want to consider the importance of a each word in the sentence. English has a lot of filler words that appear very frequently like “and”, “the”, and “a”
- When doing statistics on text, these words introduce a lot of noise since they appear way more frequently than other words
- Some NLP pipelines will flag them as **stop words** —that is, words that you might want to filter out before doing any statistical analysis
- Stop words are usually identified by just by checking a hardcoded list of known stop words
- But there’s no standard list of stop words that is appropriate for all applications. The list of words to ignore can vary depending on your application

Dependency Parsing

- The next step is to figure out how all the words in our sentence relate to each other. This is called *dependency parsing*.
- The goal is to build a tree that assigns a single **parent** word to each word in the sentence
- The root of the tree will be the main verb in the sentence

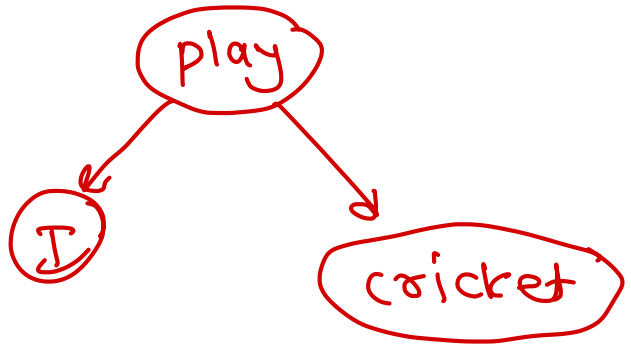

Finding Noun Phrases

- So far, we've treated every word in our sentence as a separate entity
- But sometimes it makes more sense to group together the words that represent a single idea or thing
- We can use the information from the dependency parse tree to automatically group together words that are all talking about the same thing
- Whether or not we do this step depends on our end goal
- But it's often a quick and easy way to simplify the sentence if we don't need extra detail about which words are adjectives and instead care more about extracting complete ideas.

restaurant is bad → bad

restaurant is not bad → good

Named Entity Recognition (NER)

- The goal of *Named Entity Recognition*, or *NER*, is to detect and label these nouns with the real-world concepts that they represent
- But NER systems aren't just doing a simple dictionary lookup. Instead, they are using the context of how a word appears in the sentence and a statistical model to guess which type of noun a word represents.
- Here are just some of the kinds of objects that a typical NER system can tag
  - People's names
  - Company names
  - Geographic locations (Both physical and political)
  - Product names
  - Dates and times
  - Amounts of money
  - Names of events

Coreference Resolution

- Coreference resolution is one of the most difficult steps in our pipeline to implement
- It's even more difficult than sentence parsing
- It is the process of finding the references from the previous sentences
  - → Discourse integration